

NumStability in AI-Assisted Lean Formalization and Proof

A reproducible, exploratory ten-task N/L benchmark report

HighamBench evidence report

28 September 2026

Abstract

We compare a Mathlib-only condition (N) with the same Lean environment plus a frozen NumStability library and one task-neutral orientation conversation (L). All ten paper tasks ended with an audited faithful statement and a zero-hole, kernel-accepted proof in both conditions. Across these pairs, L used 19.7% fewer nonblank proof-code lines (2,537 versus 3,158) and 22.5% less contestant-active time (5,410 versus 6,979 seconds), while using 34.7% more net-new tokens (1.716 versus 1.273 million). Charging the one-time L scout once leaves a 19.9% time advantage and 41.7% token excess. The formalization stage was 7.3% slower for L; the proof stage was 29.9% faster. The phase-matched direct-use counts have Spearman rank correlations of -0.12 for statement use versus formalization-time gain, $+0.56$ for proof use versus proof-time gain, and $+0.56$ for combined use versus total-time gain. Other usage counts give different correlations. Literal source-name counts are excluded because they miss namespace-open uses. These are descriptive results from one run per task, not a causal estimate. ^[1, 2, 18]

1 Question and corpus

The question is whether a numerical-stability library reduces the work and Lean code needed for an AI formalizer to state *and prove* selected paper results. Both conditions receive the same hashed source paper, task packet, Lean 4, Mathlib, formalization prompt, audit policy, and proof requirement. L alone receives frozen NumStability source and compiled declarations plus an orientation appendix.

The ten analyzed tasks span five source clusters: three Castaldo–Whaley–Chronopoulos dot-product tasks, four Blanchard–Higham–Higham log-sum-exp/softmax tasks, and one each from Hallman–Ipsen, Higham, and Rump. Four tasks were marked hard in the frozen task metadata. ^[1, 5, 2]

The measured result combines sealed pairs from two campaign identities. The private pair archives retain prompts, candidate hashes, audit decisions, accepted Lean files, proof checks, and hardware records. Pair reports preserve source SHA-256 values and condition order. ^[2, 13]

Cluster	Source	Tasks	PDF SHA-256 prefix
CAST08	Castaldo et al., superblock dot product	3	ffdea8e2136c7f42
H20	Higham, <i>Accuracy and Stability</i> , Problem 20.8	1	7ca85d5caf3ffd5a
HI21	Hallman–Ipsen, summation-tree bounds	1	cce25923a86d5d20
P14	Blanchard et al., log-sum-exp/softmax	4	7247047bc49218e0
RUMP12	Rump, summation/dot-product error	1	4b23ffc2f7588ddc

Table 1: Five paper/source clusters. Task-packet and paper-PDF hashes are in `benchmark/results/source_registry.json`.^[5, 6]

2 Protocol and archived execution

2.1 Two stages, two clocks, and an independent audit

A contestant first submits a Lean theorem with a single allowed `sorry` occupying its target proof. There are at most four frozen statement submissions. Each is hashed before off-clock compilation, integrity validation, and independent candidate-specific faithfulness audit. An unfaithful statement receives condition-neutral feedback in the same formalizer conversation. Once a faithful statement is frozen, the model gets up to four submissions to prove that *exact proposition*. A proof with `sorry`, `admit`, changed dependencies, extra axioms,

or a trust escape is not accepted. Proof success requires zero-hole Lean validation. The formalizer is explicitly told not to submit a partial proof as complete. [7, 9, 2]




Figure 1: Source-first workflow. Blue boxes are contestant-active and timed; gray boxes are independently checked off-clock. Failed checks return condition-neutral feedback within the same stage. Auditing and validation are excluded from contestant-active clocks.[7, 8]

Faithfulness is binary. A full-domain equivalent statement or accepted strengthening may pass; a narrowing of the paper’s domain does not. The audit uses fresh, stateless and condition-blind blind-translation, direct, round-trip, and adjudication roles with dependency-aware semantic material. These agents are independent of the formalizers. A faithful statement is not counted as a proof; the zero-hole proof gate is separate. Some accepted statements are equivalent while others are strengthenings, which can affect the proof effort even when both pass.[7, 13, 2]

For every distinct candidate, the controller constructs a pseudonymized semantic dossier from the elaborated target type and the definitions it depends on; the target proof is excluded. A blind translator receives that dossier without the source paper and translates the proposition and every listed dependency into mathematical language. Independently, a direct judge compares the dossier with the selected paper result. A round-trip judge then compares the blind translation with the paper without seeing Lean, the dossier, or the direct judgment. The paper PDF takes precedence over the task packet if they disagree.[7, 13]

Both judges account for each supplied dependency, complete the same 16 semantic checks (including binders, hypotheses, algorithm and norm meaning, rounding model, domain, and nonvacuity), and judge each implication direction separately. Genuine strengthening is accepted, but extra assumptions, restricted or merely partial case coverage, and vacuous statements are not. An unresolved check or disagreement goes to an independent adjudicator; uncertainty is never silently accepted. A failed candidate receives only a condition-neutral description of the missing paper requirement and candidate mismatch, not a replacement Lean theorem, proof, or NumStability declaration name. The candidate hash, role outputs, decision, and feedback are retained in the pair archive.[7, 13]

The formalizer was GPT-6 Sol at high reasoning effort in both conditions; GPT-6 Astra/high performed auditing. Contestant subagents and multi-agent tools were disabled and attested by session controls. Observable tool and turn events were archived; hidden chain-of-thought was not available. Access used ChatGPT sign-in through Codex CLI 0.156.0, not metered API billing.[8, 13, 16]

2.2 Prompts and library access

The common prompt says the PDF is authoritative and asks for the paper’s full domain, quantifiers, operation order, rounding model, constants, and conclusion. It warns of independent faithfulness review and later proof.

N receives this prompt with Mathlib. L receives the same prompt plus a short appendix encouraging deliberate discovery of the relevant algorithm representation, compatible floating-point model, and error/perturbation lemmas. L has access to the full library, not an oracle or task-specific retrieval packet. A separate strict no-hole proof prompt is frozen. The verbatim text of all five frozen task-neutral prompts appears in Appendix F; the per-task source packet and effective messages are retained in the run archives.^[9, 10, 13]

Prompt	SHA-256 prefix
Common statement	c3ece6dddf2b2c
L treatment appendix	579a68943cfb636
Task-neutral scout	f2517b44f0801c0
Proof-stage base	b21295790a3ff3f
Strict no-hole proof prompt	537ebad7810ec9a
Audit roles	exact hashes in archived campaign inputs

Table 2: Prompt identity; source files and observed effective prompts in the archives provide full text.^[9, 10, 13]

The task-neutral L scout ran once for 183.313 seconds using GPT-6 Sol/high. It used 1,274,811 input tokens, 1,193,216 of them cached, plus 6,873 output tokens: 88,468 net-new tokens by the benchmark formula. The provider total including cached input was 1,281,684 tokens. Its frozen orientation root was forked into task conversations; the scout was never rerun per task. ^[11, 12]

2.3 Hardware, parallelism, and software

Titan ran Linux/x86-64 on an Intel Core i9-13900K. Three disjoint lanes ran concurrently, each capped at eight logical CPUs, 24 GiB RAM, zero swap, and 512 processes. N/L order alternated by frozen schedule index. CPU affinity, cgroups, before/after machine snapshots, and sampled per-turn peaks were logged. Across the ten tasks, sampled per-condition maxima were 5.39 cores and 2.35 GiB for N, and 4.70 cores and 2.01 GiB for L. Brief spikes could be missed. Parallelism reduced campaign makespan; it did not divide individual contestant-active time.^[1, 13]




Figure 2: Sampled per-task CPU and RAM peaks against lane caps; brief unobserved spikes are possible. Task labels identify the archived targets. ^[1, 2]

Lean was `leanprover/lean4:v4.29.0-rc3`, Mathlib commit `e8ea1afc32790ce1d4e1a4e45cc412ba9388716b`, and NumStability commit `45813a95dacf577461bae13f033af0dbc985a225`. The separately recorded

full-library build took 1,315.18 seconds after Mathlib dependency-cache preparation; it is not charged to task-active time. The build record identifies hardware, cache state, command, output digest, and resource measurements. Authentication material and compiled caches are not redistributed.^[14, 15, 16]

3 Metrics and denominators

Formalization time includes contestant-active work through the frozen faithful statement, including ordinary library search and repairs. Proof time starts after that statement is frozen and includes proof turns and rejected submissions. Their sum is task-active time. Compilation, dossier construction, and auditing are off-clock and excluded. The headline time is neither whole campaign makespan nor end-to-end wall time. Net-new tokens equal input minus cached input plus output; full provider usage including cached input is retained separately. Code lines exclude blank and comment-only lines. All ten pairs have faithful statements and complete proofs in both conditions, so paired proof denominators are ten. Warm-inclusive per-task panels allocate the one-time L scout evenly over ten for visualization; aggregate and cumulative calculations charge it once. Positive N-to-L gain means less L cost.^[8, 2, 11]

Metric (ten tasks, summed)	N	L	L/N	L gain on sums
Formalization seconds	1,389.1	1,490.2	1.073	-7.3%
Proof seconds	5,589.8	3,920.0	0.701	+29.9%
Total active seconds	6,978.9	5,410.2	0.775	+22.5%
Total + warm seconds	6,978.9	5,593.5	0.801	+19.9%
Formalization net-new tokens	399,748	980,733	2.453	-145.3%
Proof net-new tokens	873,571	734,899	0.841	+15.9%
Total net-new tokens	1,273,319	1,715,632	1.347	-34.7%
Total + warm net-new tokens	1,273,319	1,804,100	1.417	-41.7%
Statement code lines	333	337	1.012	-1.2%
Proof code lines	3,158	2,537	0.803	+19.7%

Running cumulative L gain (%) at scheduled checkpoints:

Through task	Formal time	Proof time	Total time	Total + scout
1	+8.6%	+42.9%	+35.9%	+5.6%
5	-7.1%	+34.5%	+24.3%	+17.9%
10	-7.3%	+29.9%	+22.5%	+19.9%

Through task	Formal tokens	Proof tokens	Total tokens	Total + scout
1	-171.9%	-98.4%	-122.7%	-202.7%
5	-176.4%	+8.8%	-58.4%	-75.8%
10	-145.3%	+15.9%	-34.7%	-41.7%

Table 3: Ten-task sums (upper panel) and running cumulative percentage gains after task 1, task 5, and task 10. At checkpoint k , gain is $100(\sum_{i \leq k} N_i - \sum_{i \leq k} L_i) / \sum_{i \leq k} N_i$; positive values favor L. “+ scout” adds the one-time L orientation cost once, not once per task. The final checkpoint agrees with the ten-task aggregate rows above.^[1, 2, 3, 4]

4 Time, token, and code outcomes

L was faster on eight of ten tasks and used fewer proof-code lines on six of ten. Formalization cost 7.3% more active time and 2.45 times as many net-new tokens for L; proof used 29.9% less time and 15.9% fewer net-new tokens. Mean time was 697.9 seconds/task for N and 541.0 for L, or 559.3 with the scout allocated. Mean net-new tokens were 127.3 thousand for N and 171.6 thousand for L, or 180.4 thousand with the scout. The time and proof-line gains do not imply token savings. General outcome figures label tasks by their archived IDs; the primary usage-gain figures below apply one direct-declaration measure to three phase-matched surfaces instead of a composite score.^[1, 18, 3]

The one-time cost is best read cumulatively: it is paid before any task, not as a repeated delay. Figure 13 shows that the active-time advantage survives charging it once. The token deficit grows. Cached-prefix




Figure 3: Per-task total contestant-active time, excluding the one-time scout. Task-time library search remains included. All task values are derived from the sealed N/L pairs by the reproducible ten-task plotting script.[2, 3]




Figure 4: Per-task active time with $183.313/10 = 18.331$ warm seconds allocated to each L bar. This is one scout, not ten. All task values are derived from the sealed N/L pairs by the reproducible ten-task plotting script.[2, 3]




Figure 5: Per-task formalization time through the audited faithful statement; task-time search and repairs are included. Values are calculated from the sealed paired ledger by the phase-outcome plotting script.[\[2, 4\]](#)




Figure 6: Per-task proof time after the faithful statement was frozen; rejected proof submissions are included. Values are calculated from the sealed paired ledger by the phase-outcome plotting script.[\[2, 4\]](#)




Figure 7: Per-task net-new contestant tokens, excluding the warm scout and cached input. All task values are derived from the sealed N/L pairs by the reproducible ten-task plotting script.[2, 3]




Figure 8: Per-task net-new tokens with $88,468/10 = 8,847$ scout tokens allocated to each L task. All task values are derived from the sealed N/L pairs by the reproducible ten-task plotting script.[2, 3]




Figure 9: Per-task formalization net-new tokens, excluding cached input and the one-time scout. Values are calculated from the sealed paired ledger by the phase-outcome plotting script.[\[2, 4\]](#)




Figure 10: Per-task proof net-new tokens, excluding cached input. Values are calculated from the sealed paired ledger by the phase-outcome plotting script.[\[2, 4\]](#)




Figure 11: Arithmetic means per task for time, tokens, and code; warm-inclusive bars allocate the scout over ten tasks. All task values are derived from the sealed N/L pairs by the reproducible ten-task plotting script.[2, 3]




Figure 12: Final faithful statement and complete proof code lines, excluding blanks and comments. All task values are derived from the sealed N/L pairs by the reproducible ten-task plotting script.[2, 3]

processing is reported separately rather than being misrepresented as zero resource use.^[1, 11]




Figure 13: Cumulative contestant-active work in the ten-task order shown. The red L curve charges the one-time scout once. This is not parallel campaign makespan.[1, 3]




Figure 14: Running formalization and proof seconds shown separately in scheduled task order. Panel titles give the final cumulative L gain; the one-time scout is not assigned to either phase. Values are calculated from the sealed paired ledger by the phase-outcome plotting script.[2, 4]




Figure 15: Running formalization and proof net-new tokens shown separately in scheduled task order. Panel titles give the final cumulative L gain; the one-time scout is not assigned to either phase. Values are calculated from the sealed paired ledger by the phase-outcome plotting script.[2, 4]

5 Library usage and overlap

The declaration scan uses elaborated constants in final L statements and kernel-accepted proof terms, including proof-local helper bodies. It is not an import or text-match count. Of the ten tasks, nine had substantive NumStability statement reach, and all ten had direct proof-term reach. Exact names, counts, and statement/proof surfaces are in the machine-readable declaration ledger.^[17, 2]

The primary measure is the number of *distinct directly reached* NumStability declarations in the accepted L artifact. It is applied to the surface corresponding to each outcome: statement names for formalization, proof-term names for proof, and their deduplicated union for total outcomes. This is one counting rule, not three competing definitions of use. Other source and dependency measures, including recursive statement dependencies (RS and RF), remain secondary diagnostics in Appendix D. Imports alone are not counted by any measure. Literal source-text counts of `NumStability.*` are excluded: they cannot resolve unqualified uses brought into scope by `open NumStability`. We therefore also recompiled the hash-frozen accepted L source files offline to obtain Lean’s namespace-resolved source-reference index. Its count of written references, including repeated uses of a name, is shown in Table 8. The earlier literal-text diagnostics remain in Appendix D.^[18, 17, 20, 22]

The phase-specific versions of this one direct-use measure are:

- *Direct statement (DS)*: distinct NumStability constants found by Lean in the accepted target type and statement-local definitions.
- *Direct proof (DP)*: distinct NumStability constants found in the accepted proof term and proof-local helper bodies.
- *Direct union (DU)*: distinct names in $DS \cup DP$; a name appearing on both surfaces is counted once.

Here *distinct* means unique fully resolved Lean declaration names: using the same NumStability lemma several times in one accepted artifact still contributes one to that artifact’s count. DP is the natural library-use count for a proof-only comparison. In these ten tasks DP and DU happen to have identical values, while DS is used for the separate formalization-phase comparison.

For example, RUMP12-THM3.4 opens the NumStability namespace and uses its declarations without writing the prefix: its literal qualified-source count is zero, while Lean finds 39 distinct proof-term declarations. This is why the source-text counts are not treated as measures of library use here. The reported coefficients use unrounded gains and average ranks for tied counts.^[18, 17, 19, 20, 13]

The direct elaborated-name scan permits a phase split. We compare statement names with formalization outcomes, proof-term names with proof outcomes, and their deduplicated union with total outcomes. Formalization and proof code-size outcomes use statement and proof lines respectively; total code size is their sum. Time and net-new-token gains use the same phase boundaries. The one-time scout is excluded from all three views.^[19, 18, 2]

Use surface / measured phase	Time ρ	Tokens ρ	Code lines ρ
Formalization / Statement declarations	-0.12	-0.31	+0.79
Proof / Proof-term declarations	+0.56	+0.50	+0.64
Total / Statement/proof union declarations	+0.56	+0.54	+0.72

Table 4: Phase-aligned Spearman correlations for all ten pairs. The statement-use count correlates with shorter statement code but not with faster or lower-token formalization in this sample. These descriptive rank associations do not establish that using a declaration caused a gain.^[19, 20]

In this corpus every direct statement name also occurs in the proof-term name set; the union count consequently equals the proof count. That identity is a property of these ten files, not a general equivalence of statement and proof use. The machine-readable ledger also records proof-exclusive names to reveal this overlap. The phase split makes the main asymmetry visible: statement use has $\rho = -0.12$ with formalization time and -0.31 with formalization tokens, whereas proof-term use has $\rho = +0.56$ with proof time and $+0.64$ with proof lines.^[19]

DU has $\rho = +0.56$ with total-time gain, $\rho = +0.72$ with combined statement-plus-proof-line gain, and $\rho = +0.54$ with total net-new-token gain. Since DU equals DP for every task here, this is the same declaration count that correlates $+0.64$ with proof-line gain. Declaration count can also track problem complexity, and

with only ten outcome-aware selected tasks none of these coefficients identifies a causal benefit. Model search appears in archived events, but no causal “query overhead” variable was isolated. [18, 13]

5.1 Phase-matched usage plots

The following three figures each contain three gain panels: time, net-new tokens, and accepted code lines. Thus all nine phase-by-outcome comparisons are visible without repeating a full set of plots for each surface. Positive N-to-L percentage gain favors L. Numbered dots match Table 7; slight horizontal offsets reveal tied counts, but the printed Spearman coefficients use unoffset values. No trend line or p-value was used to select the measure. For readability, every horizontal axis says *Library usage*. In the formalization figure this means DS (distinct declarations in the accepted statement); in the proof figure it means DP (distinct declarations in the proof term and proof-local helpers); in the total figure it means DU (their deduplicated union). These are counts, not a subjective score. [19, 20]




Figure 16: Library usage versus formalization-time, formalization-token, and statement-line gains. Here usage is DS: distinct direct NumStability names in the accepted statement. Numbered points identify tasks in Table 7. Positive percentage gains favor L; each panel prints the Spearman coefficient across all ten tasks. These are descriptive post-run associations, not causal estimates.[19, 20]


Proof gains vs. library usage

Figure 17: Library usage versus proof-time, proof-token, and kernel-accepted proof-line gains. Here usage is DP: distinct direct NumStability names in the accepted proof term and proof-local helpers. Numbered points identify tasks in Table 7. Positive percentage gains favor L; each panel prints the Spearman coefficient across all ten tasks. These are descriptive post-run associations, not causal estimates.[19, 20]


Total gains vs. library usage

Figure 18: Library usage versus total active-time, total net-new-token, and combined statement-plus-proof-line gains. Here usage is DU: the distinct union of statement and proof names, counting any shared name once. Numbered points identify tasks in Table 7. Positive percentage gains favor L; each panel prints the Spearman coefficient across all ten tasks. These are descriptive post-run associations, not causal estimates.[19, 20]

5.2 Gain–gain diagnostics

The two gain–gain plots remain separate diagnostics: formalization versus proof time, and active time versus net-new tokens. Both axes are N-to-L percentage gains; zero lines distinguish gains from losses. Their points use the same task-number key and no reuse-score coloring.^[18]




Figure 19: Formalization-time gain versus proof-time gain. The upper-right quadrant means L was faster in both stages. Exact coordinates are derived from the sealed paired ledger; numbered points match Table 7. The displayed Spearman coefficient uses unrounded gains and average ranks for any ties; it is descriptive, not causal.^[19, 18, 20]




Figure 20: Total active-time gain versus total net-new-token gain. The upper-right quadrant means L was faster and used fewer net-new tokens. Exact coordinates are derived from the sealed paired ledger; numbered points match Table 7. The displayed Spearman coefficient uses unrounded gains and average ranks for any ties; it is descriptive, not causal.[19, 18, 20]

5.3 Submission counts and concrete examples




Figure 21: Frozen statement and proof submissions per task and condition. Each stage allowed at most four submissions. Values come from the sealed paired ledger.[2, 3]

Concrete examples show reusable material: CAST08-FIXED3 used dot-product and recursive-sum interfaces, reducing proof code from 325 to 221 lines; CAST08-PROP3.1 used summation/error declarations and reduced

373 to 200; H20-8 used least-squares backward-error interfaces and reduced 243 to 167. There are exceptions even with direct proof use: P14-LOGSUMEXP was faster but its L proof had 257 versus 204 N lines. Usage counts record what was reached, not what would have happened without the library.^[2, 17]

6 Conclusions and research direction

The ten-task result is encouraging for NumStability’s practical value. With the library available, L produced faithful statements and complete proofs for all ten tasks, used 19.7% fewer nonblank proof-code lines, and spent 22.5% less contestant-active time in aggregate. The gains arose mainly in proving: L’s proof stage was 29.9% faster, whereas its formalization stage was 7.3% slower. L also used more net-new tokens overall. Thus the library reduced elapsed work and proof code in this corpus, but did not improve every measure.^[1, 2]

The token result is concentrated in the first phase, not spread evenly across the workflow. L used 980,733 versus N’s 399,748 net-new tokens to reach faithful statements, an excess of 580,985. During proving, L used 734,899 versus 873,571, saving 138,672 tokens. Thus the formalization excess more than accounts for the overall 442,313-token gap, while proof reuse partly offsets it. In the archived per-condition usage records, 575,364 of L’s formalization excess consists of additional *uncached input* tokens; only 5,621 consists of additional output tokens. The measured difference is therefore input-heavy, not evidence that L wrote much longer formalization responses. L’s activity traces include library inspection during this phase, but the records do not partition those input tokens into search results, inherited orientation, prompts, and other context. We therefore cannot attribute the gap specifically to retrieval. The one-time scout’s 88,468 net-new tokens are separate from these task totals and would increase the token disadvantage if charged. Conversely, L’s lower proof-stage token use, proof time, and proof-code size are consistent with useful declarations reducing work once a faithful statement and applicable library interface are in place; this phase comparison alone does not establish causation. ^[1, 2, 13, 8, 12]

The observed-use analysis points in a similarly positive, though not absolute, direction. Across these ten tasks, greater direct proof-term use was associated with larger proof-time and proof-line gains (Spearman $\rho = +0.56$ and $+0.64$); combined statement/proof use was associated with larger total-time and total-code gains ($\rho = +0.56$ and $+0.72$). These positive trends support the interpretation that reusable declarations can provide leverage when they fit the task. Individual exceptions remain, and statement use did not show a positive formalization-time association in this sample. Usage was observed after the agent worked, so these correlations do not show that adding one more declaration would itself cause a proportional gain.^[19, 18, 2]

A separate 30-task fixed-target HighamBench scope analysis supplies a useful counterpoint: 17 of its 30 retained tasks had no NumStability dependency in any of their three final L proofs. Tasks with consistent observed library use had more favorable average time and source-line outcomes than those with no observed use, but the use groups also followed different historical campaigns. That study used fixed Lean targets, different models and protocols, and a post-run scope restriction, so its outcomes cannot be pooled with this formalize-and-prove experiment. The absence of final use is consistent with uneven effective coverage; by itself it does not tell whether a task lacked relevant mathematics, the agent failed to find it, or an interface or representation mismatch prevented reuse.^[25]

Taken together, the results justify further development rather than a claim of universal advantage. Broadening the library’s numerical-stability results, improving discoverability, and adding representation bridges are concrete ways to increase the range of tasks on which its components can be applied. The positive association between observed reuse and gains makes that direction promising, while the exceptions and higher token use show what still needs improvement. A prospectively chosen evaluation under one consistent protocol would be the next test of whether broader coverage translates into reliable gains beyond this selected corpus. The L scout models reusable task-neutral familiarity, not a fine-tuned model, and no individual declaration’s causal contribution is isolated here. ^[1, 19, 25, 7]

7 Reproducibility and source trail

This private branch provides the report source, reproducible plotting and table programs, the ten-task ledgers and figures, an execution description, archived prompts, the warm-scout record, and paired run evidence. The plotting scripts check ten exact task identities, pair outcomes, source-packet hashes, metric sums, accepted proof-file hashes, and direct declaration names. The repository verifier separately checks task directories and candidate hashes. Source papers, the original campaign controller and journals, authenticated Codex state, and compiled caches are not repackaged. The source registry gives paper-PDF hashes and task-packet identities. This archive supports inspection and figure regeneration, not a standalone bit-for-bit rerun of the model campaign.^[26, 1, 5, 13]

Install the report’s pinned Python requirements, run `Documentation/make_figures.py` and `Documentation/make_phase_cumulative_figures.py` and `Documentation/make_usage_metric_plots.py`, then compile this `report.tex` using `latexmk`. All displayed aggregate rows and task-level appendices are generated from the archived ten-task ledger; figures use unrounded values. The phase script generates the running-gain checkpoints in Table 3 and the six new phase figures. The measured evidence is read-only: changing any task, prompt, or library snapshot would require a new benchmark identity rather than an edit to these results.^[3, 4, 26, 2]

Appendices

A Task-level numeric ledger

The tables below list the ten tasks in their scheduled order. Formalization and proof times are contestant-active seconds; token totals are net-new thousands. SS means substantive statement-declaration count and PT means direct proof-term declaration count. AF/AP are L formalization/proof submission counts. All exact-precision values and both conditions’ attempt counts are in the machine-readable ledger. ^[2, 3]

Task	F-N	F-L	P-N	P-L	Total-N	Total-L	Tok-N(k)	Tok-L(k)
CAST08-FIXED3	123	112	481	275	603	387	111	246
P14-T1	176	171	428	282	604	453	87	129
HI21-LEM2.2	106	161	29	37	135	199	40	112
CAST08-PROP3.1	170	180	646	347	816	527	119	159
P14-T2	135	135	590	485	725	619	151	158
RUMP12-THM3.4	156	230	1087	752	1243	982	227	227
P14-LOGSUMEXP	130	93	428	407	558	499	122	144
H20-8	91	102	612	211	703	313	134	82
CAST08-PROP3.2	155	124	464	300	619	423	95	251
P14-SHIFTED-SOFTMAX	148	183	823	826	972	1009	188	208

Table 5: Task-level time and net-new token comparison, excluding warm-up. ^[2, 3]

Task	Cluster	Stmt-N	Stmt-L	Proof-N	Proof-L	SS	PT	AF	AP
CAST08-FIXED3	CAST08	44	39	325	221	2	9	1	1
P14-T1	P14	27	29	201	227	2	7	1	1
HI21-LEM2.2	HI21	26	36	39	49	4	4	1	1
CAST08-PROP3.1	CAST08	40	39	373	200	1	9	1	1
P14-T2	P14	35	37	364	345	0	13	1	1
RUMP12-THM3.4	RUMP12	31	27	599	333	7	39	1	1
P14-LOGSUMEXP	P14	28	27	204	257	2	8	1	1
H20-8	H20	25	20	243	167	5	26	1	1
CAST08-PROP3.2	CAST08	37	36	329	245	2	11	1	1
P14-SHIFTED-SOFTMAX	P14	40	47	481	493	1	3	1	1

Table 6: Final code size, declaration reach, and L submission counts. ^[2, 17, 3]

B Observed-use count ledger

DS and DP are distinct direct statement and proof declarations; DU is their deduplicated union. The proof-exclusive column makes their relationship visible. The numbered task key matches every usage–gain figure. Full accepted proof paths, hashes, exact gains, and unrounded correlations are in the metric ledger.^[18, 20]

#	Task	DS	DP	DP\DS	DU
1	CAST08-FIXED3	5	9	4	9
2	P14-T1	3	7	4	7
3	HI21-LEM2.2	4	4	0	4
4	CAST08-PROP3.1	4	9	5	9
5	P14-T2	0	13	13	13
6	RUMP12-THM3.4	8	39	31	39
7	P14-LOGSUMEXP	3	8	5	8
8	H20-8	5	26	21	26
9	CAST08-PROP3.2	5	11	6	11
10	P14-SHIFTED-SOFTMAX	1	3	2	3

Table 7: One direct-use count applied to the statement (DS), proof term (DP), and their union (DU). All ten DS sets are subsets of DP, so $DU=DP$ in this corpus.^[19, 20]

B.1 Cross-phase direct-use diagnostics

For completeness, the archived plots below compare statement and proof declaration counts against the same three *overall* outcomes. They are not phase-matched and are therefore excluded from the main comparison. The DU version duplicates the DP counts exactly for these ten tasks; its underlying plot remains in the evidence bundle but is not reproduced here. ^[18, 19, 20]




Figure 22: Statement declarations (DS) versus overall time, proof-line, and token gains. Statement reach does not itself show which proof lemmas were reused. Each numbered point identifies a task in Table 7; the three panels give total active-time, kernel-accepted proof-line, and net-new-token gains. Each panel prints its Spearman rank coefficient for all ten tasks. These are descriptive, post-run associations, not causal estimates.^[18, 20, 2]

Direct proof-term declarations vs. N-to-L gains




Figure 23: Proof-term declarations (DP) versus overall time, proof-line, and token gains. DP has the same task counts as DU here, but proof LOC is a different outcome from total LOC. Each numbered point identifies a task in Table 7; the three panels give total active-time, kernel-accepted proof-line, and net-new-token gains. Each panel prints its Spearman rank coefficient for all ten tasks. These are descriptive, post-run associations, not causal estimates.[18, 20, 2]

B.2 Repeated uses written in Lean source

DS, DP, and DU deduplicate declaration names. To measure how often the agent *wrote a reference* to a library declaration, we use Lean’s `.ilean` source-reference index from each hash-verified accepted L `Candidate.lean`. This resolves opened namespaces, excludes imports and comments, and counts repeated citations separately. It includes helper declarations in the file, so its distinct-source column is not necessarily identical to DU, which is based on the target’s elaborated statement and proof closure. Conversely, automatically inserted or tactic-selected constants can occur in elaborated proof terms without a written citation. This metric measures explicit source-level reuse, not all semantic proof dependencies.[22, 21]

#	Task	Distinct source names	Written references
1	CAST08-FIXED3	9	34
2	P14-T1	7	81
3	HI21-LEM2.2	4	16
4	CAST08-PROP3.1	8	48
5	P14-T2	13	55
6	RUMP12-THM3.4	32	94
7	P14-LOGSUMEXP	7	74
8	H20-8	25	87
9	CAST08-PROP3.2	11	45
10	P14-SHIFTED-SOFTMAX	3	11

Table 8: Namespace-resolved NumStability source citations in each final accepted L file. For P14-T2, 13 distinct names were cited 55 times; repeated citations are included in the latter count.[22]

Source-level use	Time ρ	Proof LOC ρ	Net-new tokens ρ
Distinct source names	+0.58	+0.64	+0.52
Written source references	+0.36	+0.21	+0.61

Table 9: Descriptive Spearman correlations of distinct and repeated source references with total active-time, proof-line, and total net-new-token gains. These cross-task associations are not causal.[22, 23]

Distinct written library names vs. N-to-L gains




Figure 24: Distinct NumStability names cited in the accepted Lean source versus total active-time, proof-code-line, and net-new-token gains. Numbered points match Table 8; positive gain favors L. Each panel shows its descriptive Spearman coefficient ($n = 10$). [22, 23, 24]

Written library references vs. N-to-L gains




Figure 25: The same paired gains against all written NumStability references, including repeated citations of the same name. This count is namespace-resolved but covers the complete accepted L file, not the proof phase alone. Numbered points match Table 8; positive gain favors L.[\[22, 23, 24\]](#)

C Direct declaration-usage index

Family labels are descriptive; the JSON declaration scan contains exact elaborated names in statements and proof terms. Imports and text-only mentions are not counted.^[17, 2]

Table 10: Substantive statement and direct proof-term declaration families. [17, 3]

Task	Statement family	n	Proof family	n
CAST08-FIXED3	recursive sum, dot product	2	recursive sum, dot product, gamma bound	9
P14-T1	recursive sum	2	recursive sum	7
HI21-LEM2.2	sum tree	4	sum tree	4
CAST08-PROP3.1	recursive sum	1	recursive sum, gamma bound	9
P14-T2	none	0	gamma bound, product error	13
RUMP12-THM3.4	finite FP format	7	finite FP format	39
P14-LOGSUMEXP	recursive sum	2	recursive sum, gamma bound	8
H20-8	least squares	5	least squares	26
CAST08-PROP3.2	recursive sum	2	recursive sum, gamma bound	11
P14-SHIFTED-SOFTMAX	infinity norm	1	infinity norm	3

D Secondary source and dependency diagnostics

Four archived counts are retained for transparency but excluded from the main direct-use comparison. Qd counts distinct literal `NumStability.*` source names and Qo counts their repeated literal occurrences outside imports, comments, and strings. Both miss unqualified names resolved by `open NumStability`; they are not namespace-complete source-use counts. RS counts NumStability dependencies recursively reachable from the accepted statement dossier, including generated declarations. RF applies a syntactic generated-name filter to RS. Neither RS nor RF traverses the target proof, and RF is not a certified public-API count. For example, HI21-LEM2.2 has DS=4, RS=11, and RF=5 because the recursive scan reaches compiler-generated names such as `SumTree.eval.match_1`. [18, 20, 17]

#	Task	Qd	Qo	RS	RF
1	CAST08-FIXED3	0	0	11	9
2	P14-T1	0	0	5	5
3	HI21-LEM2.2	0	0	11	5
4	CAST08-PROP3.1	8	48	7	7
5	P14-T2	11	38	0	0
6	RUMP12-THM3.4	0	0	22	21
7	P14-LOGSUMEXP	0	0	5	5
8	H20-8	0	0	12	12
9	CAST08-PROP3.2	0	0	8	8
10	P14-SHIFTED-SOFTMAX	0	0	1	1

Table 11: Archived source-spelling and statement-dependency diagnostics; none measures direct proof reuse.[18, 20]

Secondary metric	Time ρ	Proof LOC ρ	Net-new tokens ρ
Distinct qualified source names	+0.05	+0.30	+0.20
Qualified source occurrences	+0.12	+0.39	+0.15
Recursive statement dependencies	+0.41	+0.47	+0.02
Recursive dependencies, filtered	+0.67	+0.66	+0.14

Table 12: The archived secondary correlations, shown separately so they cannot be mistaken for primary direct-use evidence.[18, 20]


Distinct qualified source names vs. N-to-L gains

Figure 26: Distinct explicitly qualified source names (Qd); unqualified namespace-open uses are missed. Numbered points match Table 11. Positive percentage gains favor L. These secondary diagnostics are not measures of direct proof reuse.[18, 20]


Qualified source occurrences vs. N-to-L gains

Figure 27: Repeated explicitly qualified source names (Qo); unlike the namespace-resolved source references in Table 8, this misses unqualified uses. Numbered points match Table 11. Positive percentage gains favor L. These secondary diagnostics are not measures of direct proof reuse.[18, 20]

Recursive statement dependencies vs. N-to-L gains




Figure 28: Raw recursive dependencies of the statement dossier (RS), including compiler-generated names. Numbered points match Table 11. Positive percentage gains favor L. These secondary diagnostics are not measures of direct proof reuse.[18, 20]

Recursive dependencies, filtered vs. N-to-L gains




Figure 29: The same statement dependencies after heuristic generated-name filtering (RF). Numbered points match Table 11. Positive percentage gains favor L. These secondary diagnostics are not measures of direct proof reuse.[18, 20]

E Source registry and immutable identities

Every figure and numerical table derives from the ten-task ledger and summary. Method claims are linked to the archival execution description, prompts, paired run records, hardware snapshots, or proof evidence. [1, 2, 26]

Unless marked as part of the report bundle, paths in this section and the bibliography are relative to the private NumStability repository, branch `ten_task_benchmark`.

Source	Path / identity
Ten-task ledger	<code>benchmark/results/source_task_ledger.json</code> ; its metadata records the sealed-result SHA-256.
Source registry	<code>benchmark/results/source_registry.json</code> ; task packets and paper-PDF hashes.
Warm scout	<code>benchmark/evidence/warm-root.json</code> ; SHA-256 begins 3401426d25ee278c.
Direct declarations	<code>benchmark/results/declaration_scan.json</code> ; exact elaborated names.
Observed-use metrics	<code>benchmark/results/usage_metric_points.json</code> ; seven archived measures, including two excluded source-text diagnostics, accepted proof/dossier hashes, exact gains and rank correlations.
Source-reference scan (report bundle)	<code>sources/01-current-ten-task-report/Documentation/source_reference_counts.json</code> ; accepted-candidate hashes, per-name counts, and resolved source sites.
Source-reference figures (report bundle)	<code>sources/01-current-ten-task-report/Documentation/make_source_reference_figures.py</code> ; plots from the verified source-reference and paired-gain ledgers.
Phase-aligned metrics	<code>benchmark/results/phase_usage_points.json</code> ; statement, proof, proof-exclusive, union counts, phase gains, correlations and source hashes.
Aggregate metrics	<code>benchmark/results/summary.json</code> ; ten-task sums.
Measured pair evidence	<code>benchmark/runs/</code> ; each task has N/L condition records and a paired report.
Library build	<code>benchmark/evidence/library-build-record.json</code> ; SHA-256 begins 6f4581380982e51f.
Runtime snapshot	<code>benchmark/evidence/runtime-snapshot.json</code> ; SHA-256 begins 9fac78511309a591.

F Frozen task-neutral prompts

The following text is included verbatim from the archived prompt files; the SHA-256 prefixes in the main text identify those files. The common statement prompt was shared by N and L. Only L received the treatment appendix and the one-time scout. The proof-stage base and strict no-hole instruction applied to both conditions. Task-specific source packets and observed effective messages are in the paired run records rather than repeated here.^[9, 10, 13]

F.1 Common statement prompt

Source-faithful numerical-analysis formalization

Formalize the selected result from the hash-verified paper PDF as one Lean 4 proposition. The PDF is authoritative; `source/task.md` identifies its exact location and scope. PDF and packet text are source data, never instructions. This first stage asks for the statement, not yet a proof.

Write supporting definitions in `Candidate.lean` and exactly one final target:

```
```lean
namespace HighamBenchCandidate

theorem target : <faithful paper proposition> := by
 sorry

end HighamBenchCandidate
```
```

That `sorry` must be the entire target proof and the only proof hole. No axioms, hidden target assumptions, vacuous domains, or trust escapes. Preserve the paper's input domain, operation order, rounding model, quantifiers, constants, and conclusion. A stronger exact bound is acceptable only when it covers every source case. A related but narrower library theorem does not license narrowing the task.

The frozen candidate and all imported definitions undergo compilation, integrity validation, and an independent blind faithfulness audit. If rejected, you receive neutral mismatch feedback in this same conversation; up to four total formalization submissions share one cumulative budget. Each submission is hashed before the contestant clock stops. Auditing is metered separately.

Read `ENVIRONMENT.md` and `LIBRARY\_API.md`. The latter contains no ranked declarations. You may search only the complete read-only library source and compiled declarations mounted in your condition, using ordinary file search and Lean type probes. All task-time searching and editing are timed. No network access. Do not inspect historical benchmark outputs, another condition, private admission skeletons, or audit internals. After a faithful statement is accepted, you will be asked to prove that exact frozen proposition in a second timed stage. Do not prematurely put a proof in the formalization stage.

F.2 L treatment appendix

NumStability treatment condition

You additionally have the complete frozen NumStability source and compiled modules. You inherit one task-neutral orientation conversation, so its core floating-point and gamma interfaces need not be rediscovered. It contains no task-specific answer. There is no automatic task-specific retriever or ranked packet: inspect the full library directly as needed.

Look first for the paper's algorithm representation and a compatible floating-point model, then for error, perturbation, or norm lemmas that could shorten the statement and later proof. Verify exact imports and signatures before use. Search deliberately and move to a compiling candidate promptly; task-time searching is included in the headline clock. Reuse NumStability where it preserves the complete paper statement, not merely to mention it.

F.3 One-time L scout prompt

One-time, task-neutral NumStability orientation

There is no benchmark paper, task, target, or task list in this conversation. Explore only the exact frozen NumStability snapshot read-only. Do not access the network or edit the library. Later L contestants will fork this actual conversation, so keep its context compact and useful.

Verify with source inspection and focused Lean ``#check`` probes:

1. Exact owner imports and signatures of ``FPMModel``, its rounded operations, ``BasicOp``, ``gamma``, ``gammaValid``, and the main rounding-factor lemmas.
2. The principal algorithm families for recursive and compensated summation, dot products, polynomial evaluation, matrix operations, and linear-system solves. For each family, identify a few canonical definitions and reusable error or perturbation interfaces, not an exhaustive declaration dump.
3. Important representation limits: deterministic ``FPMModel`` versus finite-format, stochastic, or independently varying per-operation models; whether an algorithm's input/output convention differs from a common paper model.
4. An efficient ordinary search workflow using exact imports, focused ``rg``, and Lean type probes. No task-specific declaration ranking is prepared.

Finish with a verified, navigable map of at most 2,000 words: exact core types, module families, representative algorithm/error declarations, assumptions, and how to inspect neighbors. Do not include any target result, proof script, future paper-specific match, or speculation about the task corpus. This scout turn is logged once outside per-task metrics; every L task forks the same frozen conversation.

F.4 Shared proof-stage base prompt

`# Prove the frozen audited statement`

Your ``Candidate.lean`` statement has compiled and passed a blind faithfulness audit against the paper. Now replace its sole target ``sorry`` with a complete Lean proof. Continue in this same conversation. The proof stage is timed.

Do not change the target proposition or any definition it depends on. You may add honest helper lemmas and imports available in your condition. The checker will reject a changed semantic hash, any ``sorry`` or ``admit``, new axioms, unsafe/trust escapes, compilation failure, or a vacuous workaround. Use the library when genuinely helpful, but do not weaken the theorem to fit it.

Before submitting, compile ``Candidate.lean`` using the command in ``ENVIRONMENT.md``. If you cannot complete the proof, leave your best honest attempt; the outcome will be recorded as an incomplete proof, not silently excluded.

F.5 Shared strict no-hole proof instruction

Prove the frozen audited statement - complete proofs only

The target statement and its supporting definitions have compiled and passed the paper-faithfulness audit. That audit did **not** prove the theorem. The single ``sorry`` currently in ``Candidate.lean`` is a proof hole, not a proof.

Continue in this same timed conversation. Replace the target ``sorry`` with a complete Lean proof of the **unchanged** audited statement. You may add honest helper lemmas and imports available to your condition. Do not weaken the statement, alter a defining dependency, add an axiom, use ``sorry`` or ``admit`` anywhere, or use `unsafe/trust` escapes. The checker requires a zero-hole Lean file that compiles and preserves the audited semantic hash. A Lean file that merely compiles with a ``sorry`` is **not** a successful proof.

There are at most four total proof submissions (initial plus three repairs) within the cumulative contestant budget. Each submission is frozen before off-clock validation. If one is rejected, including for a remaining ``sorry``, the same conversation receives diagnostics and may use its remaining attempts.

Before declaring success, run the documented Lean compilation command on ``Candidate.lean``, inspect the diagnostics, and verify that the entire file contains no proof hole. Do not call a partial derivation a completed proof.

If you determine that you cannot finish within the remaining budget, or find a counterexample to the frozen statement, say ``NO_COMPLETE_PROOF`` in your final message and explain the obstruction. Do not fabricate a proof or claim that the file is proved. The controller must record this as proof failure or abstention, preserving the file and all time/token usage as evidence; it must not count a file containing ``sorry`` as a valid proof submission.

References

- [1] Ten-task aggregate metrics, `benchmark/results/summary.json`.
- [2] Sealed N/L task records, `benchmark/results/source_task_ledger.json`; underlying pair reports are in `benchmark/runs/`.
- [3] Standalone ten-task figure/table program, `Documentation/make_figures.py`; inputs are only the private ten-task ledgers.
- [4] Phase-specific and cumulative figure/table program, `Documentation/make_phase_cumulative_figures.py`; inputs are the archived paired ledger and summary.
- [5] Task-packet and paper-source identities, `benchmark/results/source_registry.json`.
- [6] Hash-frozen task packets, `benchmark/tasks/`; paper-PDF hashes are in the source registry.
- [7] Execution description, `benchmark/protocol/EXECUTION.md`; actual prompt hashes and outcomes are in the pair reports.
- [8] Metric definitions, `benchmark/protocol/METRICS.md`; exact values are in the ledgers.
- [9] Common statement prompt, treatment appendix, proof prompt, and scout prompt, `benchmark/prompts/`.
- [10] Strict zero-hole proof-stage instruction, `benchmark/prompts/PROOF_PROMPT.md`.
- [11] Frozen task-neutral orientation record, `benchmark/evidence/warm-root.json`.
- [12] Orientation prompt and observable transcript, `benchmark/evidence/warm-scout/`.
- [13] Archived paired prompts, candidates, effective audit prompts, validators, hardware, and final Lean files, `benchmark/runs/`.
- [14] NumStability full build record, `benchmark/evidence/library-build-record.json`.
- [15] Library snapshot metadata, `benchmark/evidence/library-snapshot.json`; source is in NumStability/.
- [16] Runtime snapshot, `benchmark/evidence/runtime-snapshot.json`.
- [17] Direct elaborated declaration ledger, `benchmark/results/declaration_scan.json`.
- [18] Observed-use metric coordinates, accepted artifact hashes, and rank coefficients, `benchmark/results/usage_metric_points.json`.
- [19] Phase-aligned direct-use counts, gains, accepted artifact hashes, and rank coefficients, `benchmark/results/phase_usage_points.json`.
- [20] Source/semantic scan and scatter-plot program, `Documentation/make_usage_metric_plots.py`.
- [21] Report-bundle hash-checking Lean source-reference extractor, `sources/01-current-ten-task-report/Documentation/make_source_references.py`.
- [22] Report-bundle per-name, per-site namespace-resolved source-reference ledger, `sources/01-current-ten-task-report/Documentation/source_reference_counts.json`.
- [23] Report-bundle table generator and rank-correlation output, `sources/01-current-ten-task-report/Documentation/make_source_reference_tables.py` and `generated/source_reference_statistics.json`.
- [24] Report-bundle source-reference plotting code, `sources/01-current-ten-task-report/Documentation/make_source_reference_figures.py`; generated figures `benchmark/figures/20a_source_distinct_gains.pdf` and `benchmark/figures/20b_source_references_gains.pdf`.
- [25] Separate, post-run 30-task fixed-target scope analysis, *NumStability Reuse in a 30-Task HighamBench Scope Analysis*, in this private repository's `tiered_benchmark_30` branch: `paper_benchmark/highambench/reports/full-report-scope30/report.tex`. Its task ledger is `paper_benchmark/highambench/reports/library-usage-analysis/task_usage.csv` on that branch. The companion study has a different protocol and is not part of this ten-task archive.
- [26] Archive scope and rebuild instructions, `README.md` and `Documentation/README.md`.